

Dealing with thrashing

- Approach 1: working set
 - Thrashing viewed from a caching perspective: given locality of reference, how big a cache does the process need?
 - Or: how much memory does process need in order to make reasonable progress (its working set)?
 - Only run processes whose memory requirements can be satisfied
- Approach 2: page fault frequency (PFF)
 - Thrashing viewed as poor ratio of “page fetch” to “useful work”
 - $PFF = \text{page faults} / \text{instructions executed}$
 - If PFF rises above threshold, process needs more memory. If not enough memory on the system, swap out.
 - If PFF sinks below threshold, memory can be taken away

Two-level scheduler

- Divide processes into active and inactive
 - Active – means working set resident in memory
 - Inactive – working set intentionally not loaded
- Balance set: union of all active working sets
 - Must keep balance set smaller than physical memory
- Use long-term scheduler
 - Moves processes from active to inactive state until balance set is small enough
 - Periodically allows inactive to become active
 - As working set changes, must update balance set
- Complications
 - How to choose idle time threshold T ?
 - How to pick processes for active set?
 - How to count shared memory (e.g., libc.so)?