

Virtual memory

Paging to disk

M1 MOSIG – Operating System Design

Renaud Lachaize

Acknowledgments

- Many ideas and slides in these lectures were inspired by or even borrowed from the work of others:
 - Arnaud Legrand, Noël De Palma, Sacha Krakowiak
 - David Mazières (Stanford)
 - (most slides/figures directly adapted from those of the CS140 class)
 - Remzi and Andrea Arpaci-Dusseau (U. Wisconsin)
 - Textbooks (Silberschatz et al., Tanenbaum)

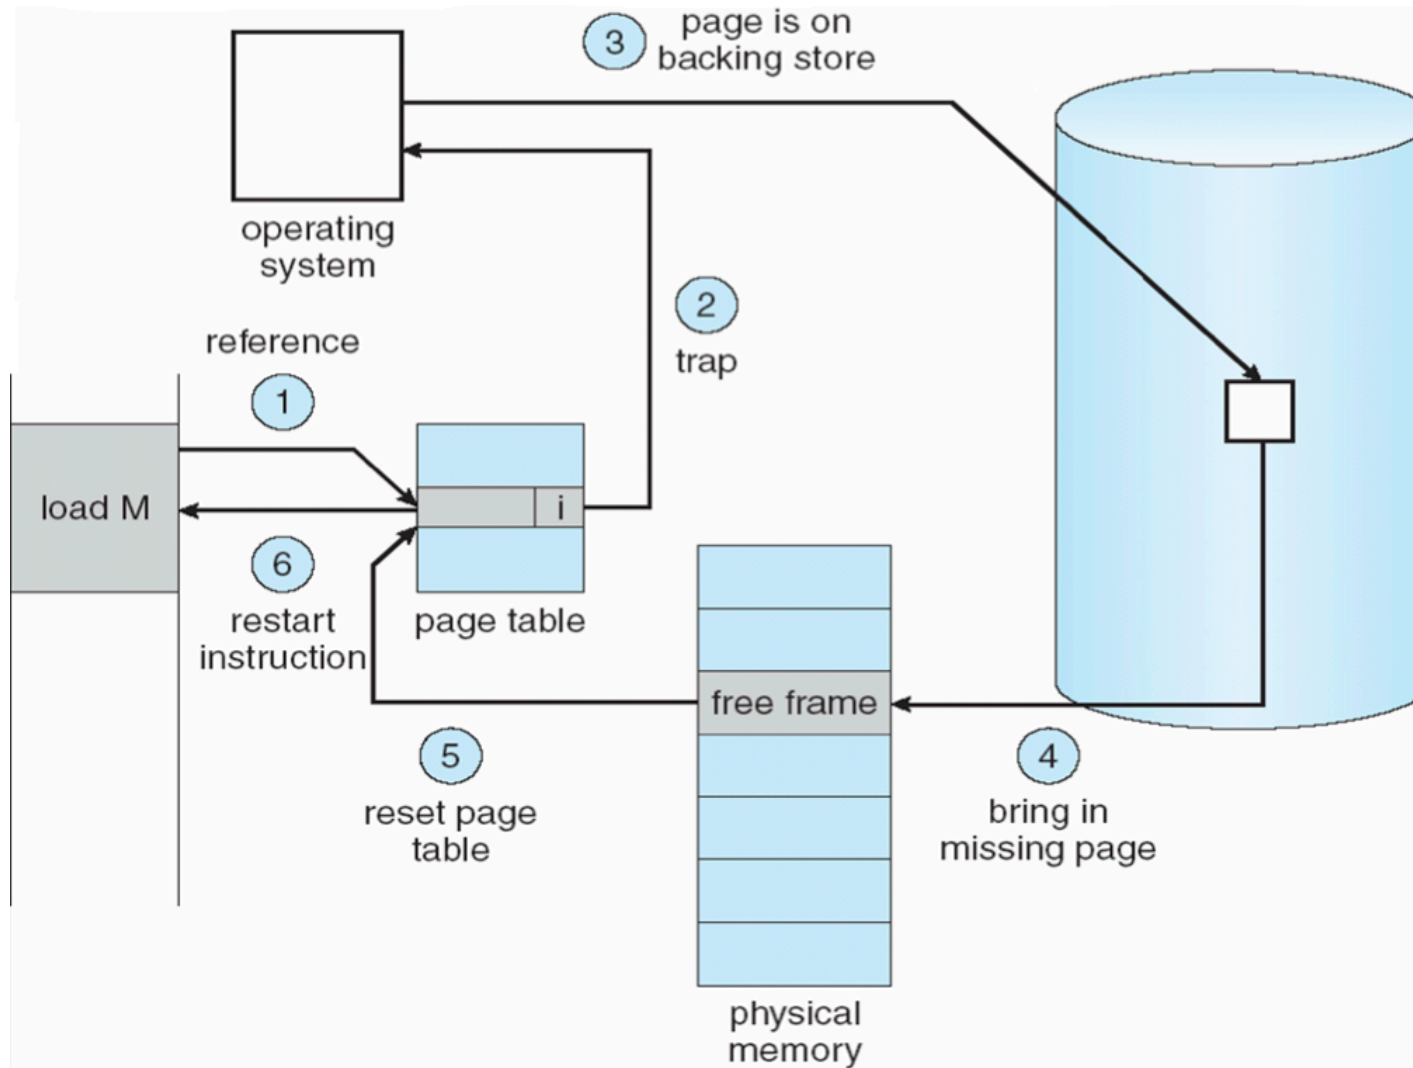
Outline

- Principles
- Challenge 1: resuming a process
- Challenge 2: choosing what to fetch
- Challenge 3: choosing what to eject
- Further problems and optimizations

Paging to disk

- Motivation: use secondary storage (disk) to provide a virtual memory with a larger capacity than the physical memory
- The RAM acts like a cache for the disk

Paging to disk (continued)




Paging to disk

Example

- **gcc** needs a new page of memory
- The kernel reclaims an idle page from **emacs**
- If the reclaimed page is **clean** (i.e., also stored on disk, with the same contents)
 - E.g., page of text from **emacs** binary on disk
 - This page can always be re-read from disk
 - OK to discard contents and give page (frame) to **gcc**
- If the reclaimed page is **dirty** (i.e., is the only valid copy)
 - The kernel must write the page to disk first before giving it to **gcc**
- Either way:
 - Mark page invalid in **emacs**'s paging information
 - **emacs** will trigger fault on next access to this virtual page
 - On fault, the kernel reads page data back from disk into a new physical page, maps new page into **emacs**, resumes execution of **emacs**

Working set model



- The disk is much, much slower than memory
 - Goal: run at memory speed, not disk speed
- 90/10 (or 80/20) rule: 10% of memory gets 90% of memory references
 - So, keep that 10% in real memory, the other 90% on disk
 - How to pick which 10%?

Paging challenges

- How to resume a process after a fault?
 - Need to save state and resume
 - Process might be in the middle of an instruction
- What to fetch?
 - Just needed page or more?
- What to evict?
 - How to allocate physical pages among processes?
 - Which pages of a particular process to keep in memory?

Re-starting instructions

- Hardware provides kernel with info about page fault
 - Faulting virtual address
 - Address of instruction that caused fault
 - Was the access a read or write? Was it an instruction fetch?
 - Was it caused by user access to kernel-only memory? (protection fault)
- Hardware must allow resuming after a fault
- Idempotent instructions are easy
 - E.g., simple load or store instruction can be restarted
 - Just re-execute any instruction that only accesses one address

What to fetch?

- Bring in page that caused page fault
- Pre-fetch surrounding pages?
 - In many cases, reading two disk blocks is approximately as fast as reading one
 - If application exhibits spatial locality, then big win to store and read multiple contiguous pages

What to evict? Selecting pages

Straw man: FIFO eviction

- Evict oldest page fetched in system
- Example - consider the following reference string:
 - 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- With a capacity of 3 physical pages: 9 page faults

1	1	4	5	9 page faults
2	2	1	3	
3	3	2	4	

What to evict? Selecting pages


Straw man: FIFO eviction

- Evict oldest page fetched in system
- Example - consider the following reference string:
 - 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- With a capacity of 3 physical pages: 9 page faults
- With a capacity of 4 physical pages: 10 page faults

1	1	5	4	
2	2	1	5	10 page faults
3	3	2		
4	4	3		

Selecting physical pages


Belady's anomaly



- More physical memory does not always mean fewer faults

Optimal page replacement

- What is optimal (if you knew the future)?
 - Replace page that will not be used for the longest period of time
- Example – with reference string
 - 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- With 4 physical pages



LRU page replacement

- Approximate optimal with least recently used
 - Because past often predicts the future
- Example – with reference string
 - 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- With 4 physical pages: 8 page faults
- Problem 1: can be pessimal – example?
 - Looping over memory (then want MRU eviction)
- Problem 2: How to implement?


1	5	
2		
3	5	4
4	3	

Straw man LRU implementations

- Idea 1: Stamp PTEs with timer value
 - E.g., using the CPU cycle counter
 - Automatically write value to PTE on each page access
 - (When page selection is needed) Scan page table to find oldest counter value = LRU page
 - Problem: would dramatically increase the memory traffic
- Idea 2: Keep doubly-linked list of pages
 - On access, remove page, place at tail of list
 - Problem: again, very expensive
- What to do?
 - Just approximate LRU, don't try to do it exactly


Clock algorithm

- Use “accessed” bit supported by most hardware
 - E.g., Pentium (x86) will write 1 to “A” bit in PTE on first access
 - Software managed TLBs like MIPS can do the same
- Do FIFO but skip accessed pages
- Keep pages in circular FIFO list
- Scan:
 - If page’s “A” bit == 1, set to 0 and skip
 - Else, if “A” == 0, evict
- A.k.a. “second-chance replacement”



Clock algorithm (continued)

- Large memory may be a problem
 - Most pages referenced in long interval
 - So we may end up having all pages with $A=1$




- Add a second clock hand
 - Two hands move in lockstep
 - Leading hand clears “A” bit
 - Trailing hand evicts pages with “A”==0
- Can also take advantage of hardware “dirty bit”
 - Each page can be (*unaccessed, clean*), (*unaccessed, dirty*), (*accessed, clean*) or (*accessed, dirty*)
 - Consider clean pages for eviction before dirty ones
- Or use n-bit variable **count** instead of just “A” bit
 - On sweep: $\text{count} = (A \ll (n-1)) \mid (\text{count} \gg 1)$
 - Evict page with lowest count

Other replacement algorithms

- Random eviction
 - Very simple to implement
 - Not overly horrible results (avoids Belady and pathological cases)
- LFU (least frequently used) eviction
 - Instead of just “A” bit, count the number of times each page is accessed
 - Least frequently accessed page must not be very useful (or maybe was just brought in and is about to be used)
 - Decay usage counts over time (for pages that fall out of usage)
- MFU (most frequently used) algorithm
 - Idea: page with the smallest count was probably just brought in and has yet to be used (so it should not be evicted)
- Neither LFU nor MFU used very commonly

Naïve paging



- Naïve page replacement: 2 disk I/Os per page fault

Page buffering

- Idea: reduce number of I/Os on the critical path
- Keep pool of free page frames
 - On fault, still select victim page to evict
 - But read fetched page into already free page
 - Can resume execution while writing out victim page
 - Then add victim page to free pool
- Can also yank pages back from free pool
 - Contains only clean pages, but may still have data
 - If page fault on page still in free pool, recycle

Outline

- Introduction to paging
 - Principle
 - Data structures and implementation examples
- Speed considerations
 - The memory wall
 - TLB
- Paging to disk
 - Principles
 - Challenge 1: resuming a process
 - Challenge 2: choosing what to fetch
 - Challenge 3: choosing what to eject
 - **Further problems and optimizations**

Page allocation

- **Allocation can be *global* or *local***
- **Global allocation does not consider page ownership**
 - E.g., with LRU, evict least recently used page of any process
 - Works well if P1 needs 10% of memory and P2 needs 70%



- Does not protect you from “memory pigs” (imagine P2 keeps looping through array that is size of mem)
- **Local allocation isolates processes (or users)**
 - Separately determine how much memory each process should have
 - Then use LRU/clock/etc. to determine which pages to evict within each process

Thrashing


- Thrashing: processes on system require more memory than it has
 - Each time one page is brought in, another page, whose contents will be soon referenced, is thrown out
 - Processes will spend all of their time blocked, waiting for pages to be fetched from disk
 - I/O devices at 100% utilization but system not getting much useful work done
- What we wanted: virtual memory as large as the disk with access time as low as the one of the physical memory
- What we have: memory with access time of the disk ☹️

Reasons for thrashing


- Process does not reuse memory, so caching does not work (past != future)
- Process does reuse memory, but it does not “fit”



- Individually, all processes fit and reuse memory, but too many for system
 - At least, this case is possible to address (see next slides)



Multiprogramming and thrashing




- Need to shed load when thrashing

Dealing with thrashing

- Approach 1: working set
 - Thrashing viewed from a caching perspective: given locality of reference, how big a cache does the process need?
 - Or: how much memory does process need in order to make reasonable progress (its working set)?
 - Only run processes whose memory requirements can be satisfied
- Approach 2: page fault frequency (PFF)
 - Thrashing viewed as poor ratio of “page fetch” to “useful work”
 - $PFF = \text{page faults} / \text{instructions executed}$
 - If PFF rises above threshold, process needs more memory. If not enough memory on the system, swap out.
 - If PFF sinks below threshold, memory can be taken away

Working sets



- Working sets changes across phases
 - Balloons during transition

Calculating the working set

- Working set: all the pages that a process will access in next T time frame
 - Cannot calculate without predicting the future
- Approximate by assuming past predicts future
 - So working set \sim pages accessed in last T interval
- Keep idle time for each page
- Periodically scan all resident pages in the system
 - “A” bit set? Clear it and clear the page’s idle time
 - “A” bit clear? Add CPU consumed since last scan to idle time
 - Working set is pages with idle time $< T$

Two-level scheduler

- Divide processes into active and inactive
 - Active – means working set resident in memory
 - Inactive – working set intentionally not loaded
- Balance set: union of all active working sets
 - Must keep balance set smaller than physical memory
- Use long-term scheduler
 - Moves processes from active to inactive state until balance set is small enough
 - Periodically allows inactive to become active
 - As working set changes, must update balance set
- Complications
 - How to choose idle time threshold T ?
 - How to pick processes for active set?
 - How to count shared memory (e.g., libc.so)?

Recap

- Paging brings nice benefits
 - Removes the fragmentation issue (in the context of address space management)
 - Enables to offload the RAM (demand paging) and thus to fit more processes in RAM
 - Enables to run processes requiring more memory than the available RAM
- Page replacement issues
 - When the RAM is full, a page must be evicted, stored back on the disk and replaced in RAM by the requested one
 - This content management has similarities with the ones in caches, TLB, ... but is implemented in software
 - Good policies build on locality, regularity of memory accesses
 - Workload and speed/size of the different memory/disk components call for different policies, data structures and tradeoffs

References

- Bruce Jacob and Trevor Mudge. Virtual memory: issues of implementation. IEEE Computer, June 1998.
- AMD and Intel documentations (see previously mentioned links)
- Replacement policies and working sets: see textbooks