

Operating Systems

Thread Synchronization

Thomas Ropars

`thomas.ropars@imag.fr`

Équipe ERODS – LIG/IM2AG/UJF

2015

Agenda

- ▶ Week 45: Synchronization: Basic Mechanisms
- ▶ Week 46: Synchronization: Implementation
- ▶ Week 47: Advanced Synchronization Techniques + IOs
- ▶ Week 48: **Second Midterm Exam** + Scheduling
- ▶ Week 49: Filesystems

References

The content of these lectures is inspired by:

- ▶ The lecture notes of Prof. André Schiper.
- ▶ The lecture notes of Prof. David Mazières.
- ▶ The lectures notes of Arnaud Legrand.
- ▶ *Operating Systems: Three Easy Pieces* by R. Arpaci-Dusseau and A. Arpaci-Dusseau

Other references:

- ▶ *Modern Operating Systems* by A. Tanenbaum
- ▶ *Operating System Concepts* by A. Silberschatz et al.

Agenda

Reminder: Threads

A Multi-Threaded Application

Mutual Exclusion

Semaphores

The Producer-Consumer Problem

Condition Variables

Monitors

More on synchronization

Agenda

Reminder: Threads

A Multi-Threaded Application

Mutual Exclusion

Semaphores

The Producer-Consumer Problem

Condition Variables

Monitors

More on synchronization

Seen previously

Threads

- ▶ Schedulable execution context
- ▶ Multi-threaded program = multiple threads in the same process address space
- ▶ Allow a process to use several CPUs
- ▶ Allow a program to overlap I/O and computation

Implementation

- ▶ Kernel-level threads
- ▶ User-level threads
 - ▶ Hybrid threading (n:m threading)
- ▶ Preemptive vs non-preemptive

POSIX threads API (pthreads)

- ▶ `tid thread_create(void (*fn)(void *), void *arg);`
- ▶ `void thread_exit();`
- ▶ `void thread_join(tid thread);`

Data sharing

- ▶ Threads share the data of the enclosing process

Concurrency issues

Threads are very useful in many contexts but using threads raises several challenges:

Concurrency issues

Threads are very useful in many contexts but using threads raises several challenges:

- ▶ How to synchronize (efficiently)?

Concurrency issues

Threads are very useful in many contexts but using threads raises several challenges:

- ▶ How to synchronize (efficiently)?
- ▶ How to cooperate (efficiently)?

Concurrency issues

Threads are very useful in many contexts but using threads raises several challenges:

- ▶ How to synchronize (efficiently)?
- ▶ How to cooperate (efficiently)?

Threads interact through shared memory:

- ▶ Memory in modern processors is very complex

Agenda

Reminder: Threads

A Multi-Threaded Application

Mutual Exclusion

Semaphores

The Producer-Consumer Problem


Condition Variables

Monitors

More on synchronization

Example: A chat server

Single-threaded version




stat. counters



users/channels

Example: A chat server

Single-threaded version




stat. counters



users/channels


Example: A chat server

Single-threaded version



Example: A chat server

First multi-threaded version




stat. counters



users/channels

Example: A chat server

First multi-threaded version




stat. counters



users/channels


Example: A chat server

First multi-threaded version




Example: A chat server

First multi-threaded version




Example: A chat server

First multi-threaded version




Example: A chat server

First multi-threaded version



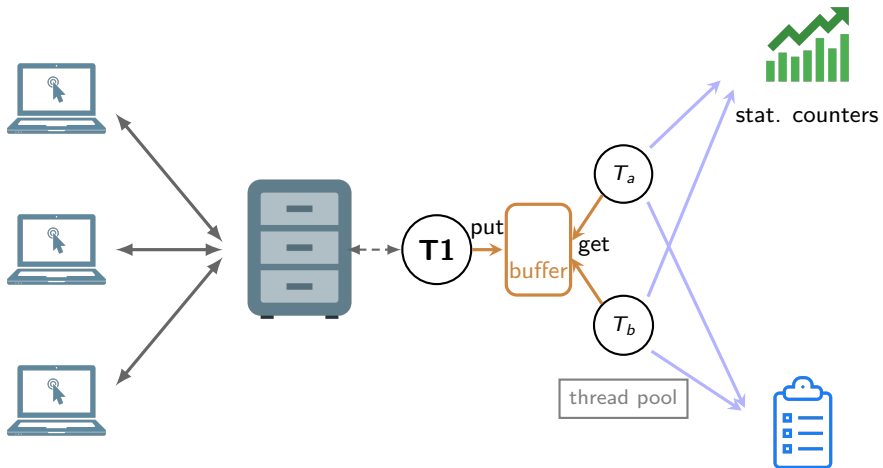
Example: A chat server

Second multi-threaded version



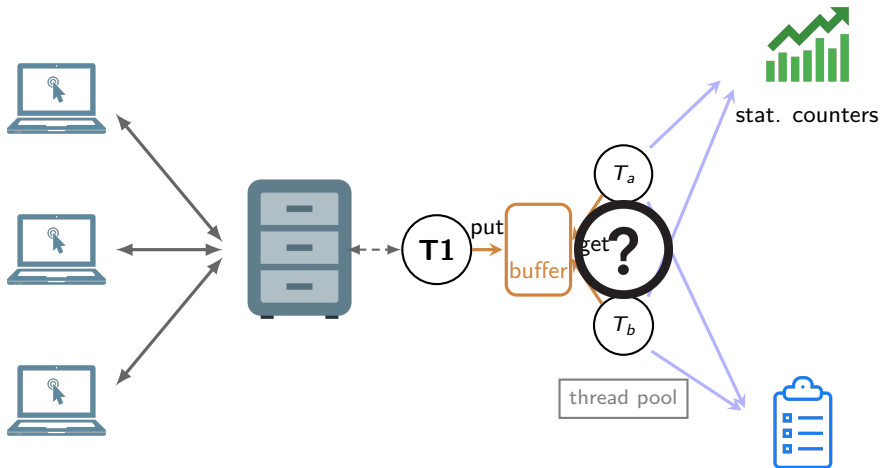
Example: A chat server

Second multi-threaded version




Example: A chat server

Second multi-threaded version



Example: A chat server

Second multi-threaded version



Classical problems

Synchronization

Mutual exclusion

- ▶ Avoid that multiple threads execute operations on the same data concurrently (*critical sections*)
- ▶ Example: Update data used for statistics

Classical problems

Synchronization

Mutual exclusion

- ▶ Avoid that multiple threads execute operations on the same data concurrently (*critical sections*)
- ▶ Example: Update data used for statistics

Reader-Writer

- ▶ Allow multiple readers or a single writer to access a data
- ▶ Example: Access to list of users and channels

Classical problems

Cooperation

Producer-Consumer

- ▶ Some threads *produce* some data that are *consumed* by other threads
- ▶ Example: A buffer of requests to be processed

Agenda

Reminder: Threads

A Multi-Threaded Application

Mutual Exclusion

Semaphores

The Producer-Consumer Problem

Condition Variables

Monitors

More on synchronization

A shared counter

```
int count = 0;
```

Thread 1:

```
for(i=0; i<10; i++){  
    count++;  
}
```

Thread 2:

```
for(i=0; i<10; i++){  
    count++;  
}
```

What is the final value of count?

A shared counter

Final value of count:

- ▶ 20
- ▶ 10
- ▶ 15
- ▶ I don't know

What do you think?

A shared counter

Final value of count:

- ▶ 20
- ▶ 10
- ▶ 15
- ▶ I don't know

A shared counter

Let's have a look at the (pseudo) assembly code for `count++`:

```
mov    count, register
add    $0x1, register
mov    register, count
```

A shared counter

Let's have a look at the (pseudo) assembly code for `count++`:

```
mov    count, register
add    $0x1, register
mov    register, count
```

A possible interleave (for one iteration on each thread)

```
mov count, register
add $0x1, register
```

```
mov register, count
```

```
mov count, register
add $0x1, register
```

```
mov register, count
```

A shared counter

Let's have a look at the (pseudo) assembly code for `count++`:

```
mov    count, register
add    $0x1, register
mov    register, count
```

A possible interleave (for one iteration on each thread)

```
mov count, register
add $0x1, register
```

```
mov register, count
```

```
mov count, register
add $0x1, register
```

```
mov register, count
```

At the end, `count=1` :-(

A shared counter

This may happen:

- ▶ When threads execute on different processor cores
- ▶ When *preemptive* threads execute on the same core

Critical section

Critical resource

A critical resource should not be accessed by multiple threads **at the same time**. It should be accessed in **mutual exclusion**.

Critical section

Critical resource

A critical resource should not be accessed by multiple threads **at the same time**. It should be accessed in **mutual exclusion**.

Critical section (CS)

A critical section is a part of a program code that access a critical resource.

Critical section: Definition of the problem

Safety

- ▶ **Mutual exclusion**: Only one thread can be in CS at a time

Liveness

- ▶ **Progress**: If no thread is currently in CS and threads are trying to access, one should eventually be able to enter the CS.
- ▶ **Bounded waiting**: Once a thread T starts trying to enter the CS, there is a bound on the number of times other threads get in.

Critical section: About liveness requirements

Liveness requirements are mandatory for a solution to be useful

Critical section: About liveness requirements

Liveness requirements are mandatory for a solution to be useful

Progress vs. Bounded waiting

- ▶ **Progress**: If no thread can enter CS, we don't have progress.
- ▶ **Bounded waiting**: If thread A is waiting to enter CS while B repeatedly leaves and re-enters C.S. ad infinitum, we don't have bounded waiting

Shared counter: New version

Thread 1:

```
Enter CS;  
count++;  
Leave CS;
```

Thread 2:

```
Enter CS;  
count++;  
Leave CS;
```

Shared counter: New version

Thread 1:

```
Enter CS;  
count++;  
Leave CS;
```

Thread 2:

```
Enter CS;  
count++;  
Leave CS;
```

How to implement `Enter CS` and `Leave CS`?

Implementation: First try using busy waiting

Shared variables:

```
int count=0;  
int busy=0;
```

Thread 1:

```
while(busy){;}  
busy=1;  
count++;  
busy=0;
```

Thread 2:

```
while(busy){;}  
busy=1;  
count++;  
busy=0;
```

Implementation: First try using busy waiting

Is the solution correct?

- ▶ The solution is correct
- ▶ The solution violates liveness
- ▶ The solution violates safety

Implementation: First try using busy waiting

Is the solution correct?

- ▶ The solution is correct
- ▶ The solution violates liveness
- ▶ The solution violates safety

Different solutions

- ▶ Peterson's algorithm
- ▶ Disable interrupts
- ▶ Locks
- ▶ Semaphores

Peterson's algorithm – Solution for 2 threads

- ▶ Assumes 2 threads T_0 and T_1

Shared variables:

```
bool wants[2] = {false, false}; /*indicates which thread  
                                wants to enter the CS*/  
int not_turn; /*not this thread turn*/
```

Code for thread T_i :

```
wants[i] = true;  
not_turn = i;  
while (wants[1-i] && not_turn == i)  
    /* other thread wants in and not our turn, so loop */;  
Critical_section ();  
wants[i] = false;  
Remainder_section ();
```


Peterson's algorithm – Correctness

Model checking

Proving the correctness of such an algorithm can be done using a model checker (e.g., UPPAAL)

Discussion about correctness

Peterson's algorithm – Correctness

Model checking

Proving the correctness of such an algorithm can be done using a model checker (e.g., UPPAAL)

Discussion about correctness

- ▶ Mutual exclusion: both threads in CS?
 - ▶ Would mean `wants[0] == wants[1] == true`,
so `not_turn` would have blocked one thread from CS

Peterson's algorithm – Correctness

Model checking

Proving the correctness of such an algorithm can be done using a model checker (e.g., UPPAAL)

Discussion about correctness

- ▶ Mutual exclusion: both threads in CS?
 - ▶ Would mean `wants[0] == wants[1] == true`,
so `no_turn` would have blocked one thread from CS
- ▶ Progress
 - ▶ If T_{1-i} doesn't want CS, `wants[1-i] == false`, so T_i won't loop
 - ▶ If both threads try to enter, one thread is the `no_turn` thread

Peterson's algorithm – Correctness

Model checking

Proving the correctness of such an algorithm can be done using a model checker (e.g., UPPAAL)

Discussion about correctness

- ▶ Mutual exclusion: both threads in CS?
 - ▶ Would mean `wants[0] == wants[1] == true`, so `not_turn` would have blocked one thread from CS
- ▶ Progress
 - ▶ If T_{1-i} doesn't want CS, `wants[1-i] == false`, so T_i won't loop
 - ▶ If both threads try to enter, one thread is the `not_turn` thread
- ▶ Bounded waiting
 - ▶ If T_i wants lock and T_{1-i} tries to re-enter, T_{1-i} will set `not_turn = 1 - i`, allowing T_i in

Peterson's algorithm – Limits

- ▶ Given solution works for 2 threads
- ▶ Can be generalized to n threads but n must be known in advance
- ▶ Note that the current version assumes that the memory is *sequentially consistent*. Most processors don't provide sequential consistency.
 - ▶ Discussed in another lecture.

Disabling interrupts

Basic idea

Prevent a thread from being interrupted while it is in CS

- ▶ On a single processor, if a thread is not interrupted, it will execute the CS atomically

Implementation

- ▶ Fully disabling interrupts is not a good solution (unsafe)
- ▶ Have a per-thread *Do-Not-Interrupt* (DNI) bit
 - ▶ Typical setup: periodic timer signal caught by thread scheduler
 - ▶ Scheduling decisions based on DNI bits

Limitations

- ▶ Cannot work on multi-processors
- ▶ Interrupts might get lost if disabled for a large amount of time

Specification

A *lock* is defined by a **lock variable** and two methods: `lock()` and `unlock()`.

Specification

A *lock* is defined by a **lock variable** and two methods: `lock()` and `unlock()`.

- ▶ A lock can be **free** or **held**

Specification

A *lock* is defined by a **lock variable** and two methods: `lock()` and `unlock()`.

- ▶ A lock can be **free** or **held**
- ▶ **lock()**: If the lock is free, the calling thread **acquires** the lock and enters the CS. Otherwise the thread is blocked.

Specification

A *lock* is defined by a **lock variable** and two methods: `lock()` and `unlock()`.

- ▶ A lock can be **free** or **held**
- ▶ **lock()**: If the lock is free, the calling thread **acquires** the lock and enters the CS. Otherwise the thread is blocked.
- ▶ **unlock()**: **Releases** the lock. It has to be called by the thread currently holding the lock.


Specification

A *lock* is defined by a **lock variable** and two methods: `lock()` and `unlock()`.

- ▶ A lock can be **free** or **held**
- ▶ **lock()**: If the lock is free, the calling thread **acquires** the lock and enters the CS. Otherwise the thread is blocked.
- ▶ **unlock()**: **Releases** the lock. It has to be called by the thread currently holding the lock.

Locks prevent programmers from having to write synchronization code explicitly.

Locks: Analogy



All global data should be protected by a lock!

- ▶ Global = accessed by more than one thread, at least one write
- ▶ Exception is initialization, before data is exposed to other threads
- ▶ This is the responsibility of the application writer

Implementation of locks

User space

- ▶ Locks can be fully implemented in user space
- ▶ Use of atomic operations + busy waiting
- ▶ More details in next lecture

With the help of the kernel

- ▶ *Block* the thread: remove it from the ready list
 - ▶ In the case of a lock, put the thread into a list of threads waiting for the lock
- ▶ *Wake up* the thread: put it back into the ready list
- ▶ *Futex* system call in Linux kernel

Pthread locks: Mutexes

- ▶ `mutex`: variable of type `pthread_mutex_t`
- ▶ `pthread_mutex_init(&mutex, ...)`: initialize the mutex
 - ▶ The macro `PTHREAD_MUTEX_INITIALIZER` can be used to initialize a mutex allocated statically with the default options
- ▶ `pthread_mutex_destroy(&mutex)`: destroy the mutex
- ▶ `pthread_mutex_lock(&mutex)`
- ▶ `pthread_mutex_unlock(&mutex)`
- ▶ `pthread_mutex_trylock(&mutex)`: is equivalent to `lock()`, except that if the mutex is held, it returns immediately with an error code

Implementation based on *Futex* in Linux

Pthread locks: Example

```
#include <pthread.h>

int count=0;
pthread_mutex_t count_mutex = PTHREAD_MUTEX_INITIALIZER;

/* ... */
pthread_mutex_lock(&count_mutex);
count++;
pthread_mutex_unlock(&count_mutex);
```


Agenda

Reminder: Threads

A Multi-Threaded Application

Mutual Exclusion

Semaphores

The Producer-Consumer Problem

Condition Variables

Monitors

More on synchronization

Semaphores

- ▶ Locks ensure mutual exclusion.
- ▶ A semaphore is another mechanism that allows controlling access to shared variables but is more powerful than a lock.
- ▶ Semaphores were proposed by Dijkstra in 1968

Semaphores

A semaphore is initialized with an integer value N and can be manipulated with two operations P and V .

Operations semantic

Semaphores

A semaphore is initialized with an integer value N and can be manipulated with two operations P and V .

Operations semantic

- ▶ P : $N = N-1$, block thread if $N < 0$
 - ▶ P stands for Proberen (Dutch) – try

Semaphores

A semaphore is initialized with an integer value N and can be manipulated with two operations P and V .

Operations semantic

- ▶ P : $N = N-1$, block thread if $N < 0$
 - ▶ P stands for Proberen (Dutch) – try
- ▶ V : $N = N+1$, wake up one thread if $N \leq 0$
 - ▶ V stands for Verhogen (Dutch) – increment

Semaphores

A semaphore is initialized with an integer value N and can be manipulated with two operations P and V .

Operations semantic

- ▶ P : $N = N-1$, block thread if $N < 0$
 - ▶ P stands for Proberen (Dutch) – try
- ▶ V : $N = N+1$, wake up one thread if $N \leq 0$
 - ▶ V stands for Verhogen (Dutch) – increment

POSIX interface

- ▶ $P \rightarrow \text{int sem_wait(sem_t *s)}$
- ▶ $V \rightarrow \text{int sem_post(sem_t *s)}$
 - ▶ Other interfaces call it `sem_signal()`

Mutual exclusion with semaphores

- ▶ Initializing a semaphore with value N can be seen as providing it with N *tokens*

Mutual exclusion with semaphores

- ▶ Initializing a semaphore with value N can be seen as providing it with N *tokens*
- ▶ To implement critical sections, a semaphore should be initialized with $N = 1$

Mutual exclusion with semaphores

- ▶ Initializing a semaphore with value N can be seen as providing it with N *tokens*
- ▶ To implement critical sections, a semaphore should be initialized with $N = 1$
 - ▶ Warning: A semaphore with $N = 1$ and a lock are not equivalent

Mutual exclusion with semaphores

- ▶ Initializing a semaphore with value N can be seen as providing it with N *tokens*
- ▶ To implement critical sections, a semaphore should be initialized with $N = 1$
 - ▶ Warning: A semaphore with $N = 1$ and a lock are not equivalent

Example

```
#include <pthread.h>

int count=0;
sem_t count_mutex;

sem_init(&count_mutex, 0, 1);
/* ... */
sem_wait(&count_mutex);
count++;
sem_post(&count_mutex);
```

Agenda

Reminder: Threads

A Multi-Threaded Application

Mutual Exclusion

Semaphores

The Producer-Consumer Problem


Condition Variables

Monitors

More on synchronization


Specification of the problem

Recall



Specification of the problem

Recall



Specification

- ▶ A buffer of fixed size
- ▶ Producer threads put items into the buffer. The *put* operation blocks if the buffer is full
- ▶ Consumer threads get items from the buffer. The *get* operation blocks if the buffer is empty

Producer Consumer

```
void producer (void *ignored) {
    for (;;) {
        /* produce an item and put in
           nextProduced */

        while (count == BUFFER_SIZE) {

            /* Do nothing */

        }

        buffer [in] = nextProduced;
        in = (in + 1) % BUFFER_SIZE;
        count++;

    }
}
```

```
void consumer (void *ignored) {
    for (;;) {

        while (count == 0) {

            /* Do nothing */

        }

        nextConsumed = buffer[out];
        out = (out + 1) % BUFFER_SIZE;
        count--;

        /* consume the item in
           nextConsumed */

    }
}
```

Producer Consumer

```
void producer (void *ignored) {  
    for (;;) {  
        /* produce an item and put in  
         nextProduced */  
  
        while (count == BUFFER_SIZE) {  
  
            /* Do nothing */  
  
        }  
  
        buffer [in] = nextProduced;  
        in = (in + 1) % BUFFER_SIZE;  
        count++;  
  
    }  
}
```

```
void consumer (void *ignored) {  
    for (;;) {  
  
        while (count == 0) {  
  
            /* Do nothing */  
  
        }  
  
        nextConsumed = buffer[out];  
        out = (out + 1) % BUFFER_SIZE;  
        count--;  
  
        /* consume the item in  
         nextConsumed */  
  
    }  
}
```

Not correct: shared data are not protected

Producer Consumer with Locks

```
mutex_t mutex = MUTEX_INITIALIZER;

void producer (void *ignored) {
    for (;;) {
        /* produce an item and put in
           nextProduced */

        mutex_lock (&mutex);
        while (count == BUFFER_SIZE) {

            thread_yield ();

        }

        buffer [in] = nextProduced;
        in = (in + 1) % BUFFER_SIZE;
        count++;
        mutex_unlock (&mutex);
    }
}
```

```
void consumer (void *ignored) {
    for (;;) {
        mutex_lock (&mutex);
        while (count == 0) {

            thread_yield ();

        }

        nextConsumed = buffer[out];
        out = (out + 1) % BUFFER_SIZE;
        count--;
        mutex_unlock (&mutex);

        /* consume the item in
           nextConsumed */
    }
}
```


Producer Consumer with Locks

```
mutex_t mutex = MUTEX_INITIALIZER;

void producer (void *ignored) {
    for (;;) {
        /* produce an item and put in
           nextProduced */

        mutex_lock (&mutex);
        while (count == BUFFER_SIZE) {

            thread_yield ();

        }

        buffer [in] = nextProduced;
        in = (in + 1) % BUFFER_SIZE;
        count++;
        mutex_unlock (&mutex);
    }
}
```

```
void consumer (void *ignored) {
    for (;;) {
        mutex_lock (&mutex);
        while (count == 0) {

            thread_yield ();

        }

        nextConsumed = buffer[out];
        out = (out + 1) % BUFFER_SIZE;
        count--;
        mutex_unlock (&mutex);

        /* consume the item in
           nextConsumed */
    }
}
```

Not correct: If a thread enters a while loop, all threads are blocked forever (deadlock)

Producer Consumer with Locks

```
mutex_t mutex = MUTEX_INITIALIZER;
```

```
void producer (void *ignored) {  
    for (;;) {  
        /* produce an item and put in  
         nextProduced */  
  
        mutex_lock (&mutex);  
        while (count == BUFFER_SIZE) {  
            mutex_unlock (&mutex);  
            thread_yield ();  
            mutex_lock (&mutex);  
        }  
  
        buffer [in] = nextProduced;  
        in = (in + 1) % BUFFER_SIZE;  
        count++;  
        mutex_unlock (&mutex);  
    }  
}
```

```
void consumer (void *ignored) {  
    for (;;) {  
        mutex_lock (&mutex);  
        while (count == 0) {  
            mutex_unlock (&mutex);  
            thread_yield ();  
            mutex_lock (&mutex);  
        }  
  
        nextConsumed = buffer[out];  
        out = (out + 1) % BUFFER_SIZE;  
        count--;  
        mutex_unlock (&mutex);  
  
        /* consume the item in  
         nextConsumed */  
    }  
}
```

Producer Consumer with Locks

```
mutex_t mutex = MUTEX_INITIALIZER;

void producer (void *ignored) {
    for (;;) {
        /* produce an item and put in
           nextProduced */

        mutex_lock (&mutex);
        while (count == BUFFER_SIZE) {
            mutex_unlock (&mutex);
            thread_yield ();
            mutex_lock (&mutex);
        }

        buffer [in] = nextProduced;
        in = (in + 1) % BUFFER_SIZE;
        count++;
        mutex_unlock (&mutex);
    }
}
```

```
void consumer (void *ignored) {
    for (;;) {
        mutex_lock (&mutex);
        while (count == 0) {
            mutex_unlock (&mutex);
            thread_yield ();
            mutex_lock (&mutex);
        }

        nextConsumed = buffer[out];
        out = (out + 1) % BUFFER_SIZE;
        count--;
        mutex_unlock (&mutex);

        /* consume the item in
           nextConsumed */
    }
}
```

Correct ... but busy waiting (waste CPU time, slow down other threads)

Cooperation

Cooperation = Synchronization + Communication

Cooperation

Cooperation = Synchronization + Communication

- ▶ **Synchronization**: Imposing an order on the execution of instructions
- ▶ **Communication**: Exchanging information between threads

Cooperation

Cooperation = Synchronization + Communication

- ▶ **Synchronization**: Imposing an order on the execution of instructions
- ▶ **Communication**: Exchanging information between threads

Semaphores allow cooperation between threads

Producer-Consumer with semaphores

- ▶ Initialize full to 0 (block consumer on empty buffer)
- ▶ Initialize empty to N (block producer when queue full)

```
void producer (void *ignored) {  
    for (;;) {  
        /* produce an item and put in  
         nextProduced */  
  
        sem_wait(&empty);  
  
        buffer [in] = nextProduced;  
        in = (in + 1) % BUFFER_SIZE;  
        count++;  
  
        sem_post(&full)  
    }  
}
```

```
void consumer (void *ignored) {  
    for (;;) {  
        sem_wait(&full);  
  
        nextConsumed = buffer[out];  
        out = (out + 1) % BUFFER_SIZE;  
        count--;  
  
        sem_post(&empty);  
  
        /* consume the item in  
         nextConsumed */  
    }  
}
```

Producer-Consumer with semaphores

- ▶ Initialize full to 0 (block consumer on empty buffer)
- ▶ Initialize empty to N (block producer when queue full)
- ▶ An additional semaphore should be used for mutual exclusion (could use a lock instead)

```
void producer (void *ignored) {  
    for (;;) {  
        /* produce an item and put in  
         nextProduced */  
  
        sem_wait(&empty);  
        sem_wait(&mutex);  
        buffer [in] = nextProduced;  
        in = (in + 1) % BUFFER_SIZE;  
        count++;  
        sem_post(&mutex);  
        sem_post(&full);  
    }  
}
```

```
void consumer (void *ignored) {  
    for (;;) {  
        sem_wait(&full);  
        sem_wait(&mutex);  
        nextConsumed = buffer[out];  
        out = (out + 1) % BUFFER_SIZE;  
        count--;  
        sem_post(&mutex);  
        sem_post(&empty);  
  
        /* consume the item in  
         nextConsumed */  
    }  
}
```


Comments on semaphores

- ▶ Semaphores allow elegant solutions to some problems (producer-consumer, reader-writer)
- ▶ However they are quite error prone:
 - ▶ If you call `wait` instead of `post`, you'll have a deadlock
 - ▶ If you forget to protect parts of your code, you might violate mutual exclusion
 - ▶ You have “tokens” of different types, which may be hard to reason about

Comments on semaphores

- ▶ Semaphores allow elegant solutions to some problems (producer-consumer, reader-writer)
- ▶ However they are quite error prone:
 - ▶ If you call `wait` instead of `post`, you'll have a deadlock
 - ▶ If you forget to protect parts of your code, you might violate mutual exclusion
 - ▶ You have “tokens” of different types, which may be hard to reason about

This is why other constructs have been proposed

Agenda

Reminder: Threads

A Multi-Threaded Application

Mutual Exclusion

Semaphores

The Producer-Consumer Problem

Condition Variables

Monitors

More on synchronization

Condition variables (pthreads)


- ▶ A **condition variable** is a shared variable to which is associated a set of waiting threads.
- ▶ It allows a thread to explicitly put itself into a *waiting mode*.
- ▶ It is used together with a mutex: When a thread puts itself to wait, the corresponding mutex is released.
- ▶ It is often associated to a “logical condition” (hence the name)

Condition variables (pthreads)


Interface

- ▶ `cond`: variable of type `pthread_cond_t`
- ▶ `pthread_cond_init(&cond, ...)`: initialize the condition
 - ▶ The macro `PTHREAD_COND_INITIALIZER` can be used to initialize a condition variable allocated statically with the default options
- ▶ `void pthread_cond_wait(&cond, &mutex)`: atomically unlock mutex and put the thread to wait on `cond`.
- ▶ `void pthread_cond_signal(&cond)` and `pthread_cond_broadcast(&cond)`: Wake one/all the threads waiting on `cond`.

Condition variable: Analogy



Condition variable: Analogy



Producer-Consumer with condition variables

```
mutex_t mutex = MUTEX_INITIALIZER;
cond_t nonempty = COND_INITIALIZER;
cond_t nonfull = COND_INITIALIZER;

void producer (void *ignored) {
    for (;;) {
        /* produce an item and
           put in nextProduced */

        mutex_lock (&mutex);
        if (count == BUFFER_SIZE)
            cond_wait (&nonfull, &mutex);

        buffer [in] = nextProduced;
        in = (in + 1) % BUFFER_SIZE;
        count++;
        cond_signal (&nonempty);
        mutex_unlock (&mutex);
    }
}
```

```
void consumer (void *ignored) {
    for (;;) {
        mutex_lock (&mutex);
        if (count == 0)
            cond_wait (&nonempty, &mutex);

        nextConsumed = buffer[out];
        out = (out + 1) % BUFFER_SIZE;
        count--;
        cond_signal (&nonfull);
        mutex_unlock (&mutex);

        /* consume the item
           in nextConsumed */
    }
}
```

Does it work with many readers and many writers?

Producer-Consumer with condition variables

```
mutex_t mutex = MUTEX_INITIALIZER;
cond_t nonempty = COND_INITIALIZER;
cond_t nonfull = COND_INITIALIZER;

void producer (void *ignored) {
    for (;;) {
        /* produce an item and
           put in nextProduced */

        mutex_lock (&mutex);
        while (count == BUFFER_SIZE)
            cond_wait (&nonfull, &mutex);

        buffer [in] = nextProduced;
        in = (in + 1) % BUFFER_SIZE;
        count++;
        cond_signal (&nonempty);
        mutex_unlock (&mutex);
    }
}
```

```
void consumer (void *ignored) {
    for (;;) {
        mutex_lock (&mutex);
        while (count == 0)
            cond_wait (&nonempty, &mutex);

        nextConsumed = buffer[out];
        out = (out + 1) % BUFFER_SIZE;
        count--;
        cond_signal (&nonfull);
        mutex_unlock (&mutex);

        /* consume the item
           in nextConsumed */
    }
}
```

Always put a while around the waiting on a condition!

Producer-Consumer with condition variables

```
mutex_t mutex = MUTEX_INITIALIZER;
cond_t nonempty = COND_INITIALIZER;
cond_t nonfull = COND_INITIALIZER;

void producer (void *ignored) {
    for (;;) {
        /* produce an item and
           put in nextProduced */

        mutex_lock (&mutex);
        while (count == BUFFER_SIZE)
            cond_wait (&nonfull, &mutex);

        buffer [in] = nextProduced;
        in = (in + 1) % BUFFER_SIZE;
        count++;
        cond_signal (&nonempty);
        mutex_unlock (&mutex);
    }
}
```

```
void consumer (void *ignored) {
    for (;;) {
        mutex_lock (&mutex);
        while (count == 0)
            cond_wait (&nonempty, &mutex);

        nextConsumed = buffer[out];
        out = (out + 1) % BUFFER_SIZE;
        count--;
        cond_signal (&nonfull);
        mutex_unlock (&mutex);

        /* consume the item
           in nextConsumed */
    }
}
```

Beware: this solution does not warrant First Come First Served!

More on conditions variables

Why must `cond_wait` both release mutex and sleep? Why not separate mutexes and condition variables?

```
while (count == BUFFER_SIZE) {  
    mutex_unlock (&mutex);  
    cond_wait(&nonfull);  
    mutex_lock (&mutex);  
}
```

More on conditions variables

Why must `cond_wait` both release mutex and sleep? Why not separate mutexes and condition variables?

```
while (count == BUFFER_SIZE) {  
    mutex_unlock (&mutex);  
    cond_wait(&nonfull);  
    mutex_lock (&mutex);  
}
```

A thread could end up stuck waiting because of a bad interleaving

```
PRODUCER  
while (count == BUFFER_SIZE){  
    mutex_unlock (&mutex);  
  
    cond_wait (&nonfull);  
}
```

```
CONSUMER  
  
mutex_lock (&mutex);  
...  
count--;  
cond_signal (&nonfull);
```

Agenda

Reminder: Threads

A Multi-Threaded Application

Mutual Exclusion

Semaphores

The Producer-Consumer Problem

Condition Variables

Monitors

More on synchronization

Monitors

- ▶ A monitor is a synchronization construct
- ▶ It provides synchronization mechanisms similar to mutex + condition variables. (Some people call both “monitors”)

Definition

- ▶ A monitor is an object/module with a set of methods.
- ▶ Each method is executed in mutual exclusion
- ▶ Condition variables (or simply “*conditions*”) are defined with the same semantic as defined previously

Comments on monitors

- ▶ Proposed by Brinch Hansen (1973) and Hoare (1974)
- ▶ Possibly less error prone than raw mutexes
- ▶ Basic synchronization mechanism in Java
- ▶ Different *flavors* depending on the semantic of signal:
 - ▶ Hoare-style: The signaled thread get immediately access to the monitor. The signaling thread waits until the signaled threads leaves the monitor.
 - ▶ Mesa-style (java): The signaling thread stays in the monitor.
- ▶ Semaphores can be implemented using monitors and monitors can be implemented using semaphores

Agenda

Reminder: Threads

A Multi-Threaded Application

Mutual Exclusion

Semaphores

The Producer-Consumer Problem

Condition Variables

Monitors

More on synchronization

The Reader-Writer problem

Problem statement

- ▶ Several threads try to access the same shared data, some reading, other writing.
- ▶ Either a single writer or multiple readers can access the shared data at any time

Different flavors

- ▶ Priority to readers
- ▶ Priority to writers

The Dining Philosophers problem

Proposed by Dijkstra

Problem statement

5 philosophers spend their live alternatively thinking and eating. They sit around a circular table. The table has a big plate of rice but only 5 chopsticks, placed between each pair of philosophers. When a philosopher wants to eat, he has to peak the chopsticks on his left and on his right. Two philosophers can't use a chopstick at the same time. How to ensure that no philosopher will starve?

Goals

- ▶ Avoid **deadlocks**: Each philosopher holds one chopstick
- ▶ Avoid **starvation**: Some philosophers never eat